

First Part - Spark

Here we run some simple queries, using Spark.

Figure 1: Answer for first query.

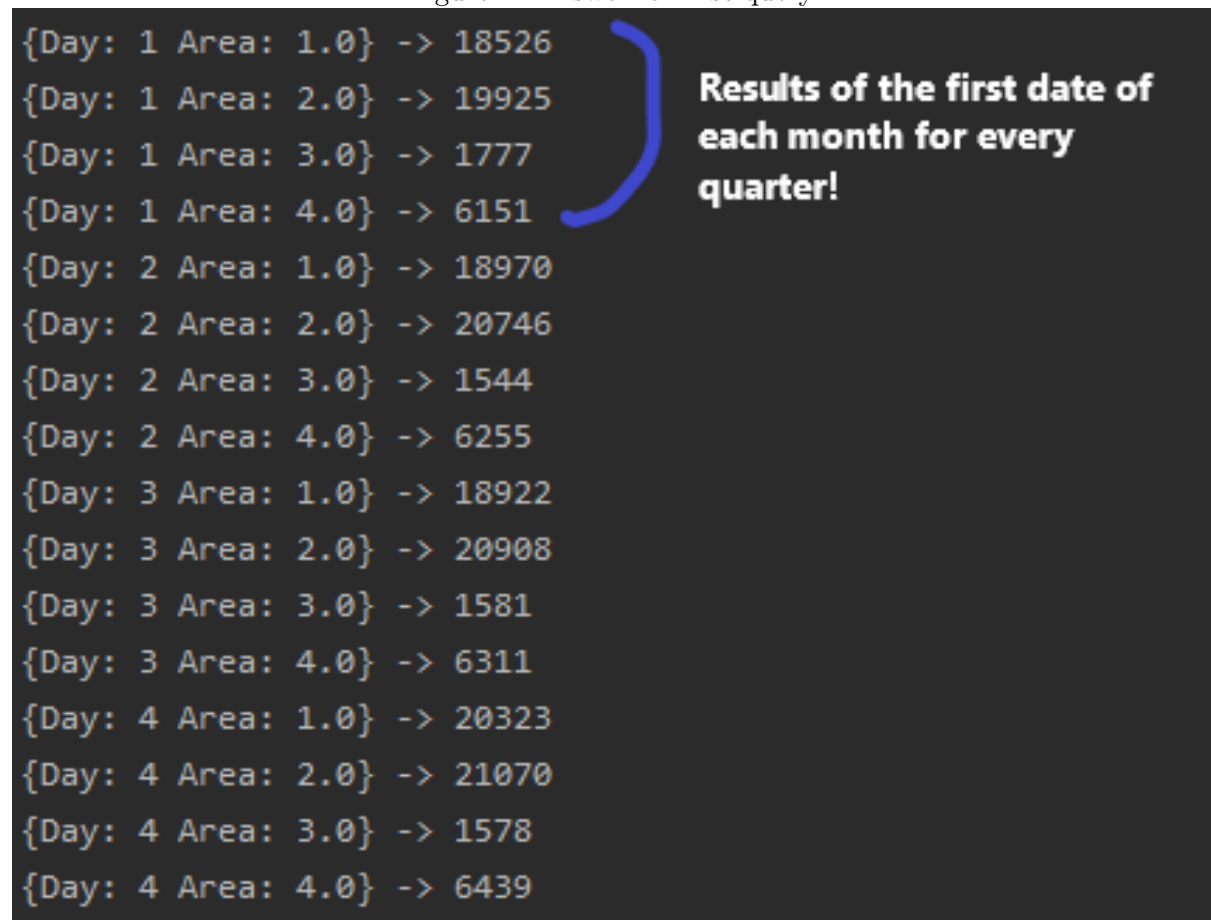


Figure 2: Answer for second query.

```
-----Second Query-----
```

District_start	trip_duration
1.0	909.4345134540057
4.0	1016.599808211573
3.0	2136.2074441991663
2.0	896.9452866314485

District_start
3.0

Figure 3: Answer for second query.

District_start	average_distance
1.0	2.720747557556261
4.0	3.134196381255205
3.0	12.093144468972284
2.0	2.3466148927027777

District_start
3.0

Figure 4: Answer for third query.

Distance	trip_duration	passenger_count
3.773	1128	4
5.068	1217	3
2.025	634	6
3.846	3528	3
2.044	721	4
2.602	870	5
5.373	1628	3
10.283	2341	3
19.859	2065	3
19.589	1884	6
1.192	962	4
1.271	773	3
2.81	714	6
5.071	909	4
2.093	1078	6
2.905	755	3
2.434	690	3
7.467	1093	5
1.975	1151	5
3.472	1030	6

Figure 5: Answer for fourth query.

```
+-----+-----+
| hour | count |
+-----+-----+
|  12  | 71522 |
|  22  | 80079 |
|   1  | 38354 |
|  13  | 71145 |
|  16  | 63982 |
|   6  | 33082 |
|   3  | 20541 |
|  20  | 83690 |
|   5  | 14911 |
|  19  | 89889 |
|  15  | 71451 |
|  17  | 76095 |
|   9  | 67383 |
|   4  | 15886 |
|   8  | 66734 |
|  23  | 69414 |
|   7  | 55347 |
|  10  | 65139 |
|  21  | 83741 |
|  11  | 68138 |
+-----+-----+
```

That shows which hours
have the most traffic!!

Figure 6: Answer for fifth query.

```
Enter longitude: 73.982154846191406
Enter latitude: 40.767936706542969
Enter hour: 12
30085
```

id vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	store
id0809232	2 2016-04-01 12:12:25	2016-04-01 12:23:17	1	-73.97953796386719	40.753360748291016	-73.96399688720703	40.763458251953125	
id1965939	1 2016-05-24 12:16:33	2016-05-24 12:28:33	1	-73.99409484863281	40.76155090332031	-73.98063659667969	40.78963851928711	
id0391524	1 2016-04-06 12:14:40	2016-04-06 12:33:18	1	-73.97864532470703	40.7480354309082	-73.97013092041016	40.76273727416992	
id1783063	2 2016-04-06 12:24:22	2016-04-06 12:40:02	2	-73.98464965820312	40.74597930908203	-73.97737884521484	40.75965118408203	
id3141299	2 2016-01-14 12:25:54	2016-01-14 12:37:24	3	-73.99577331542969	40.74256896972656	-74.00809478759766	40.72277069091797	
id1947466	2 2016-04-21 12:32:31	2016-04-21 12:37:20	6	-73.98127746582031	40.78036117553711	-73.97615051269531	40.78860092163086	
id0202431	1 2016-02-17 12:02:29	2016-02-17 12:18:44	1	-74.00413513183594	40.7428092956543	-73.98779296875	40.74486541748047	
id0363564	2 2016-05-21 12:29:39	2016-05-21 12:37:08	2	-73.99774932861328	40.756649017333984	-73.99302673339844	40.74713897705078	
id1196772	2 2016-06-28 12:08:13	2016-06-28 12:18:01	1	-73.98249053955078	40.762569427490234	-73.99862670898438	40.7606086730957	
id0238963	2 2016-01-11 12:42:36	2016-01-11 13:09:38	1	-73.98001861572266	40.78601837158203	-73.97633361816406	40.75906753540039	
id2042535	1 2016-02-12 12:36:38	2016-02-12 12:47:07	1	-73.9789810180664	40.75651931762695	-73.9921646118164	40.749202728271484	
id2975591	1 2016-01-06 12:26:54	2016-01-06 12:39:42	1	-73.99375915527344	40.73271942138672	-73.99127960205078	40.74226379394531	
id2873647	1 2016-01-18 12:54:42	2016-01-18 13:05:22	1	-73.9800033569336	40.785667419433594	-73.9644775390625	40.77593231201172	
id3152592	2 2016-02-25 12:25:30	2016-02-25 12:30:11	2	-73.98873901367188	40.75185775756836	-73.98373413085938	40.75286102294922	
id1904594	1 2016-02-04 12:25:00	2016-02-04 12:35:19	4	-73.99105834960938	40.755531311035156	-73.98719787597656	40.759254455566406	
id0461973	2 2016-02-11 12:08:38	2016-02-11 12:46:54	1	-73.98892211914062	40.73683166503906	-73.9826889038086	40.776920318603516	
id3461680	2 2016-06-01 12:03:54	2016-06-01 12:36:32	1	-73.9939193725586	40.74568176269531	-73.87319946289062	40.77436065673828	
id3834849	1 2016-04-26 12:34:46	2016-04-26 12:43:25	1	-73.99027252197266	40.7559700012207	-73.9798355102539	40.76222229003906	
id2386155	2 2016-01-24 12:42:54	2016-01-24 12:54:20	1	-73.98352813720703	40.73802947998047	-73.98323822021484	40.73794937133789	

Figure 7: Answer for sixth query.

DayOfYear	vendor_id	hour	max(count)
139	1	21	272
79	2	19	313
140	1	18	214
54	2	14	247
152	2	22	216
17	1	0	219
101	2	21	198
37	1	2	157
177	1	1	201
117	1	15	193
126	2	21	317
40	1	18	247
150	1	19	169
38	2	21	184
160	2	16	134
99	2	15	215
48	2	22	272
91	2	18	290
179	1	0	108
44	2	3	128
153	1	8	217

for every day of the year,
for every vendor, the "max"
hour!

Figure 8: Answer for seventh query.

```
-----Num_Of_Taxis in Sunday by hour-----  
  
hour -> 0:00 taxis -> 12214  
hour -> 1:00 taxis -> 10795  
hour -> 2:00 taxis -> 8548  
hour -> 3:00 taxis -> 6621  
hour -> 4:00 taxis -> 4678  
hour -> 5:00 taxis -> 1943  
hour -> 6:00 taxis -> 2132  
hour -> 7:00 taxis -> 2830  
hour -> 8:00 taxis -> 4298  
hour -> 9:00 taxis -> 6330  
hour -> 10:00 taxis -> 8661  
hour -> 11:00 taxis -> 10074  
hour -> 12:00 taxis -> 10589  
hour -> 13:00 taxis -> 10576  
hour -> 14:00 taxis -> 10545  
hour -> 15:00 taxis -> 10136  
hour -> 16:00 taxis -> 9762  
hour -> 17:00 taxis -> 10635  
hour -> 18:00 taxis -> 11059  
hour -> 19:00 taxis -> 9997  
hour -> 20:00 taxis -> 9092  
hour -> 21:00 taxis -> 8682  
hour -> 22:00 taxis -> 7845  
hour -> 23:00 taxis -> 6332
```

Figure 9: Answer for seventh query.

```
----Num_Of_Taxis in Saturday by hour----

hour: 0:00 taxis -> 11702
hour: 1:00 taxis -> 9936
hour: 2:00 taxis -> 8192
hour: 3:00 taxis -> 5997
hour: 4:00 taxis -> 3810
hour: 5:00 taxis -> 1987
hour: 6:00 taxis -> 2519
hour: 7:00 taxis -> 3699
hour: 8:00 taxis -> 5458
hour: 9:00 taxis -> 7843
hour: 10:00 taxis -> 9363
hour: 11:00 taxis -> 10227
hour: 12:00 taxis -> 11187
hour: 13:00 taxis -> 11321
hour: 14:00 taxis -> 10822
hour: 15:00 taxis -> 10818
hour: 16:00 taxis -> 9961
hour: 17:00 taxis -> 11206
hour: 18:00 taxis -> 12853
hour: 19:00 taxis -> 12992
hour: 20:00 taxis -> 11172
hour: 21:00 taxis -> 11333
hour: 22:00 taxis -> 12317
hour: 23:00 taxis -> 13120
```


Second Part - Spark-Hadoop-Docker

Here we made two datanodes and upload the data file.

Figure 10:

192.168.99.100:9870/dfshealth.html#tab-overview




Summary

Security is off.
Safemode is off.
12 files and directories, 2 blocks (2 replicated blocks, 0 erasure coded block groups) = 14 total filesystem object(s).
Heap Memory used 20.05 MB of 40.24 MB Heap Memory. Max Heap Memory is 239.75 MB.
Non Heap Memory used 62.57 MB of 63.69 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.


Configured Capacity:	35.69 GB
Configured Remote Capacity:	0 B
DFS Used:	385.64 MB (1.06%)
Non DFS Used:	4.64 GB
DFS Remaining:	28.79 GB (80.66%)
Block Pool Used:	385.64 MB (1.06%)
DataNodes usages% (Min/Median/Max/stdDev):	1.06% / 1.06% / 1.06% / 0.00%
Live Nodes	2 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	2
Number of Blocks Pending Deletion (including replicas)	0
Block Deletion Start Time	Fri Jan 10 22:51:54 +0200 2020
Last Checkpoint Time	Fri Jan 10 22:51:55 +0200 2020

Figure 11: The data file.

Browse Directory

/user/root/input Go!   

Show entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	root	supergroup	191.3 MB	Jan 10 23:24	3	128 MB	fares.csv	

Showing 1 to 1 of 1 entries

Previous **1** Next

Hadoop, 2019.

Figure 12:

NameNode Journal Status

Current transaction ID: 32	
Journal Manager	State
FileJournalManager(root=/hadoop/dfs/name)	EditLogFileOutputStream(/hadoop/dfs/name/current/edits_inprogress_000000000000000009)

NameNode Storage

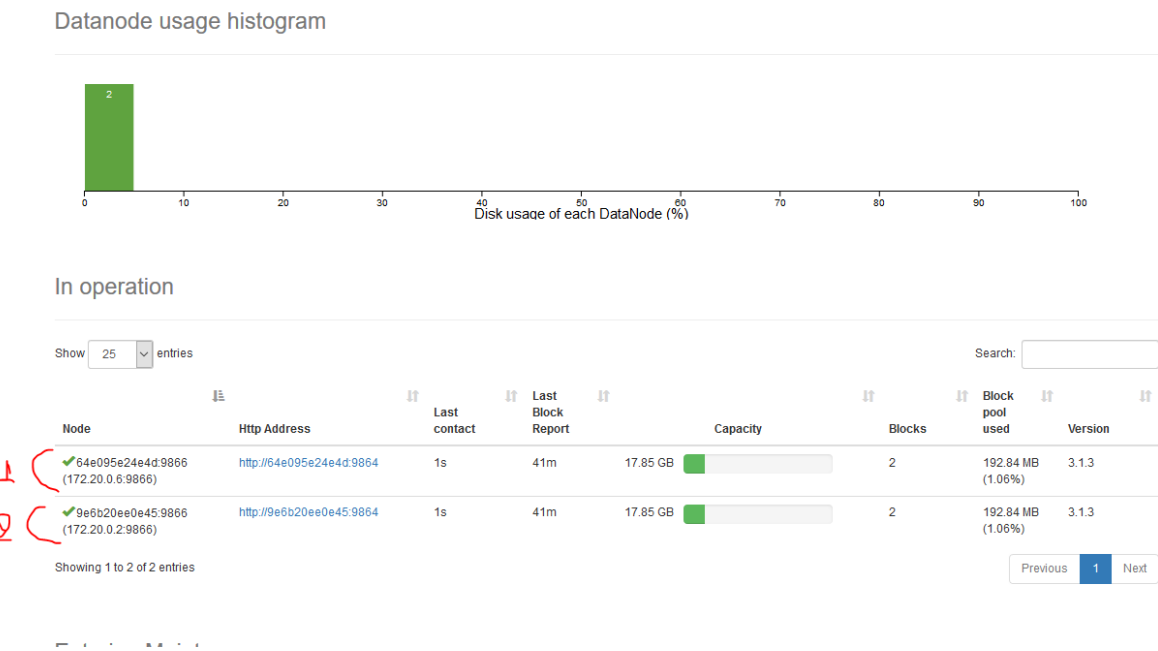
Storage Directory	Type	State
/hadoop/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	35.69 GB	385.64 MB (1.06%)	28.79 GB (80.66%)	385.64 MB	2

Hadoop, 2019.

Figure 13: The two nodes.



Run the previous queries again!

Figure 14: The two nodes.

```
root@ec10bb671ddf:/spark/bin# ./spark-submit --class testspark --master local[*] SparkHadoop-1.0-SNAPSHOT.jar
2020-01-10 22:00:44,878 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... u

Type which query you want to execute
1
-----First Query-----
root
|-- Day: integer (nullable = true)
|-- District_start: double (nullable = true)
|-- count: long (nullable = false)

{Day: 1 Area: 1.0} -> 18526
{Day: 1 Area: 2.0} -> 19925
{Day: 1 Area: 3.0} -> 1777
{Day: 1 Area: 4.0} -> 6149
{Day: 2 Area: 1.0} -> 18967
{Day: 2 Area: 2.0} -> 20746
{Day: 2 Area: 3.0} -> 1544
{Day: 2 Area: 4.0} -> 6255
{Day: 3 Area: 1.0} -> 18922
{Day: 3 Area: 2.0} -> 20908
{Day: 3 Area: 3.0} -> 1581
{Day: 3 Area: 4.0} -> 6311
{Day: 4 Area: 1.0} -> 20321
{Day: 4 Area: 2.0} -> 21070
{Day: 4 Area: 3.0} -> 1580
{Day: 4 Area: 4.0} -> 6439
{Day: 5 Area: 1.0} -> 19603
{Day: 5 Area: 2.0} -> 21703
{Day: 5 Area: 3.0} -> 1693
{Day: 5 Area: 4.0} -> 6932
{Day: 6 Area: 1.0} -> 19379
{Day: 6 Area: 2.0} -> 21356
{Day: 6 Area: 3.0} -> 1581
{Day: 6 Area: 4.0} -> 6903
{Day: 7 Area: 1.0} -> 19281
{Day: 7 Area: 2.0} -> 20807
{Day: 7 Area: 3.0} -> 1527
{Day: 7 Area: 4.0} -> 6831
{Day: 8 Area: 1.0} -> 19309
{Day: 8 Area: 2.0} -> 20973
{Day: 8 Area: 3.0} -> 1478
{Day: 8 Area: 4.0} -> 6253
{Day: 9 Area: 1.0} -> 19325
{Day: 9 Area: 2.0} -> 21655
{Day: 9 Area: 3.0} -> 1543
{Day: 9 Area: 4.0} -> 6850
{Day: 10 Area: 1.0} -> 18963
```

Figure 15: The two nodes.

```
Type which query you want to execute
2
-----Second Query-----
+-----+-----+
|District_start|   trip_duration|
+-----+-----+
|           1.0| 909.4315651111585|
|           4.0|1016.6032008802168|
|           3.0|2136.2137396778676|
|           2.0| 896.9487857973359|
+-----+-----+

+-----+
|District_start|
+-----+
|           3.0|
+-----+
only showing top 1 row

+-----+-----+
|District_start|   average_distance|
+-----+-----+
|           1.0|2.7207856535269346|
|           4.0| 3.134228018553194|
|           3.0|12.093144468972284|
|           2.0|2.3466148927027777|
+-----+-----+

+-----+
|District_start|
+-----+
|           3.0|
+-----+
only showing top 1 row

Type which query you want to execute
```

Third Part - Streaming

Figure 16: The two nodes.

```
-----
Batch: 1
-----
+---+-----+
|Hour|count|
+---+-----+
| 12|   8|
| 22|  13|
|  1|   1|
| 13|  14|
|  6|   3|
| 16|   8|
|  3|   2|
| 20|  13|
|  5|   1|
| 19|  10|
| 15|  12|
| 17|  12|
|  9|   7|
|  4|   1|
|  8|  10|
| 23|  11|
|  7|   6|
| 10|   7|
| 21|  12|
| 11|  11|
+---+-----+
only showing top 20 rows

[Stage 5:=====> (192 + 8) / 200]-----
Batch: 2
-----
+---+-----+
|Hour|count|
+---+-----+
| 12|  11|
| 22|  15|
|  1|   2|
| 13|  24|
|  6|   8|
| 16|  11|
|  3|   8|
| 20|  19|
|  5|   2|
| 19|  17|
| 15|  17|
|  9|  15|
| 17|  21|
|  4|   4|
|  8|  16|
| 23|  16|
|  7|  12|
| 10|  11|
| 21|  17|
| 11|  20|
+---+-----+
only showing top 20 rows
```